



Revisiting Time-Space Tradeoffs for Function Inversion

Spencer Peters

Noah S.D.



Siyao Guo

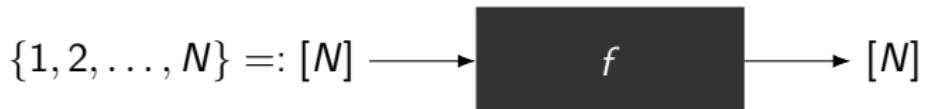


Sasha Golovnev



Function Inversion

- ▶ Given a function



Function Inversion

- ▶ Given a function, and a point y in its range,



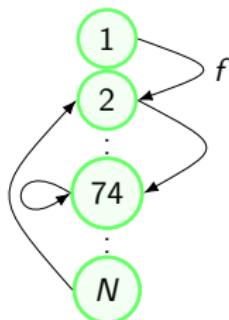
Function Inversion

- Given a function, and a point y in its range, find x with $f(x) = y$.

$$\{1, 2, \dots, N\} =: [N] \ni x \rightarrow \boxed{f} \rightarrow y \in [N]$$

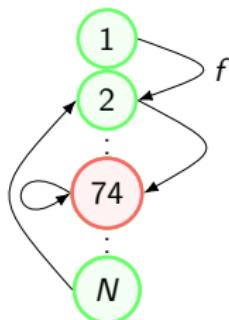
Function Inversion

- Given a function, and a point y in its range, find x with $f(x) = y$.



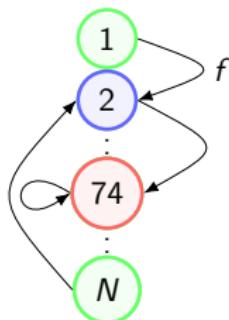
Function Inversion

- Given a function, and a point y in its range, find x with $f(x) = y$.



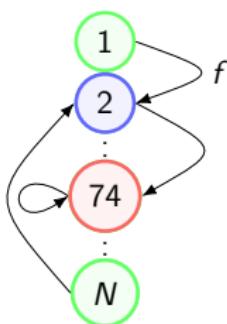
Function Inversion

- Given a function, and a point y in its range, find x with $f(x) = y$.



Function Inversion

- ▶ Given a function, and a point y in its range, find x with $f(x) = y$.



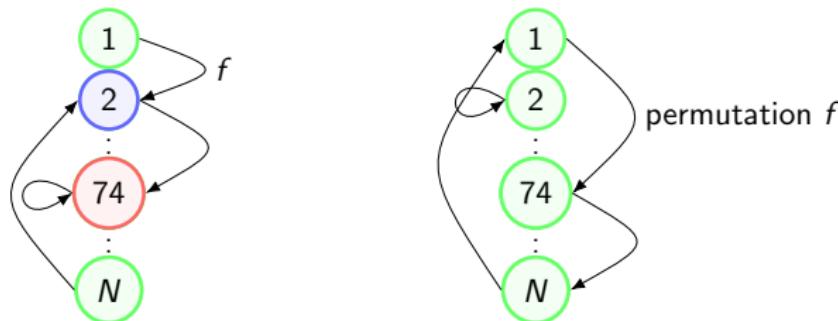
- ▶ The study of this *black-box* function inversion problem was initiated by Martin Hellman



in 1980 [Hel80].

Function Inversion

- Given a function, and a point y in its range, find x with $f(x) = y$.



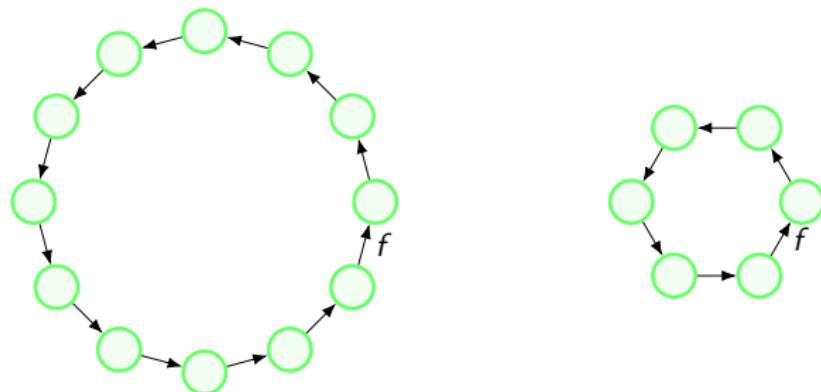
- The study of this *black-box* function inversion problem was initiated by Martin Hellman



in 1980 [Hel80].

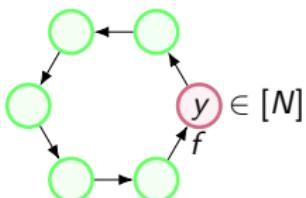
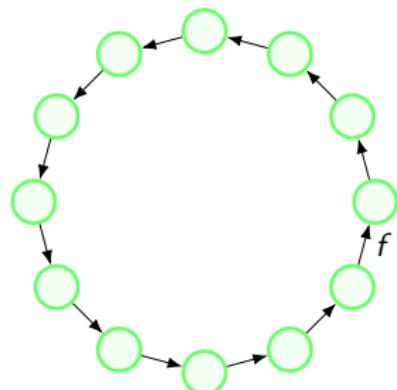
Hellman's algorithm

- If f is a permutation, its *graph* is a disjoint union of cycles.



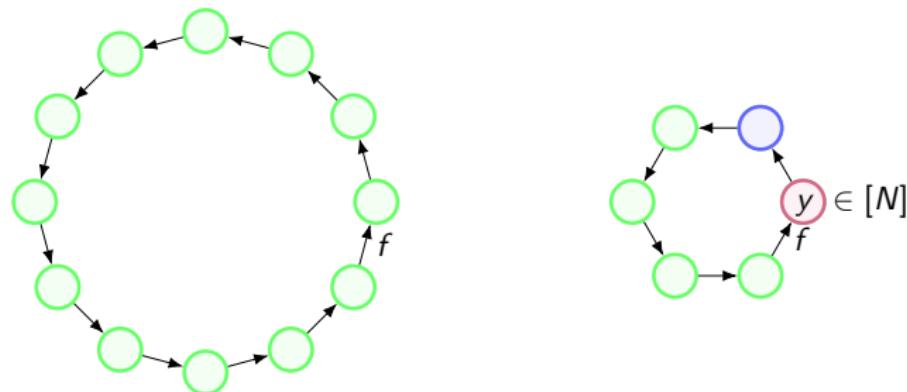
Hellman's algorithm

- If f is a permutation, its *graph* is a disjoint union of cycles.



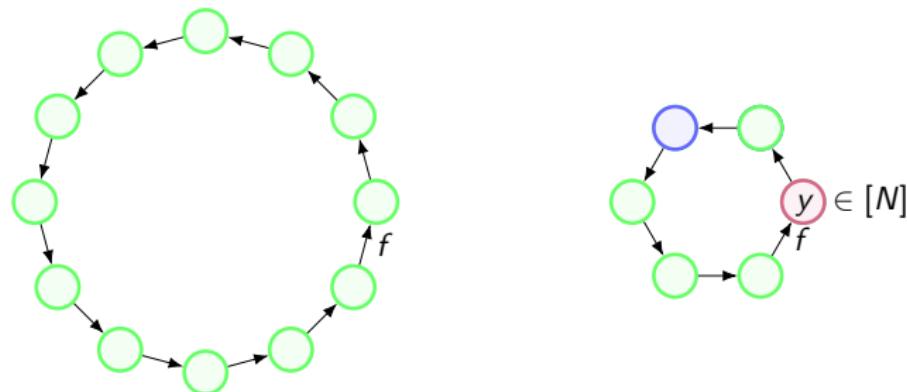
Hellman's algorithm

- If f is a permutation, its *graph* is a disjoint union of cycles.



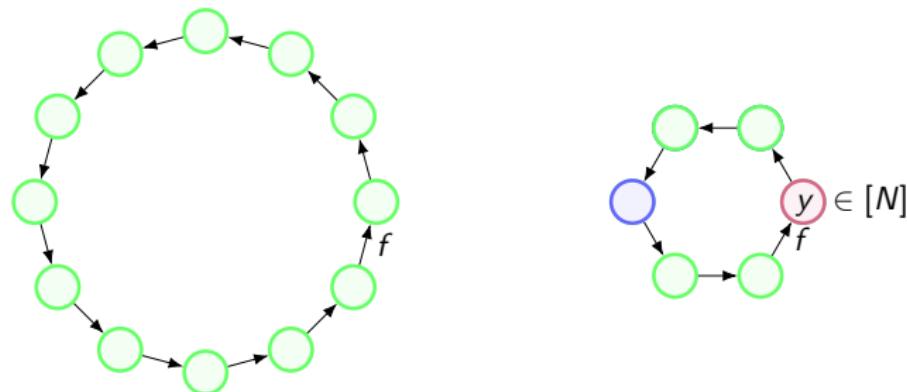
Hellman's algorithm

- If f is a permutation, its *graph* is a disjoint union of cycles.



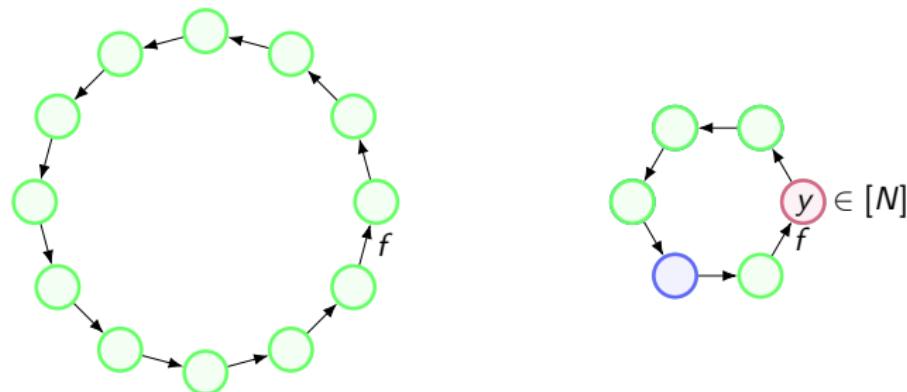
Hellman's algorithm

- If f is a permutation, its *graph* is a disjoint union of cycles.



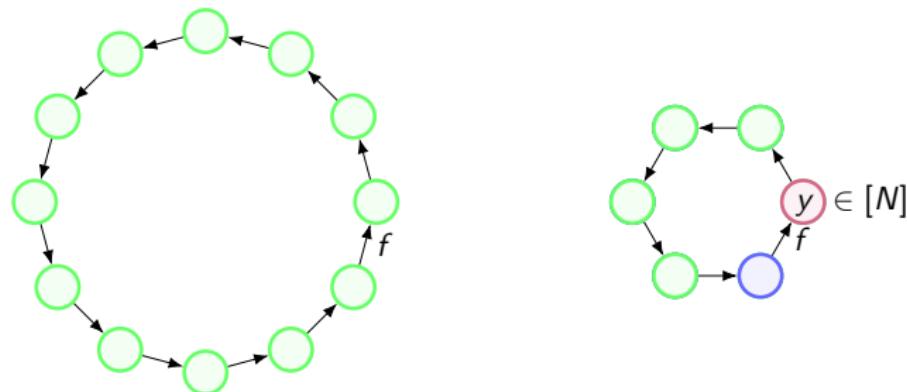
Hellman's algorithm

- If f is a permutation, its *graph* is a disjoint union of cycles.



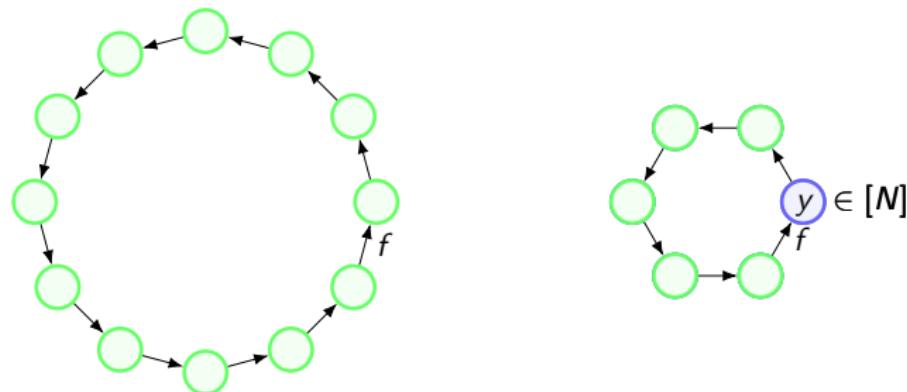
Hellman's algorithm

- If f is a permutation, its *graph* is a disjoint union of cycles.



Hellman's algorithm

- If f is a permutation, its *graph* is a disjoint union of cycles.

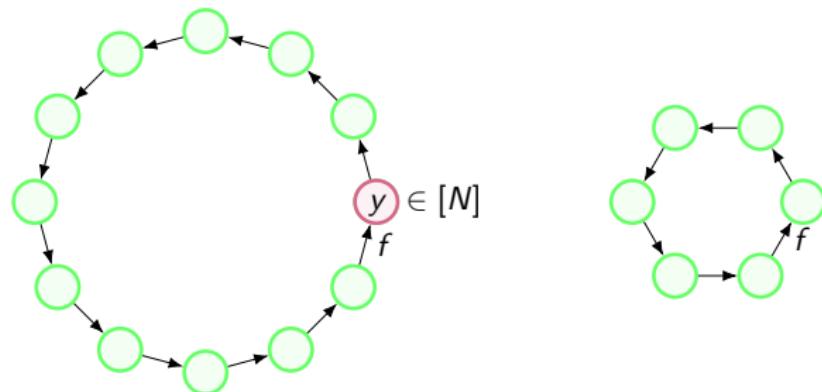


Hellman's algorithm

- ▶ What if y is on a large cycle?

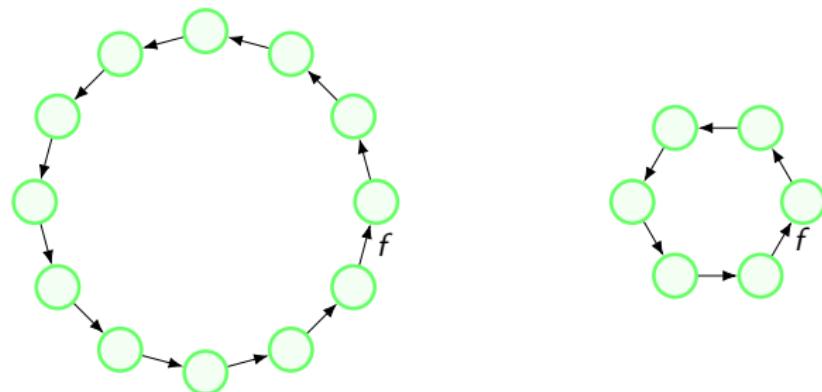
Hellman's algorithm

- ▶ What if y is on a large cycle?



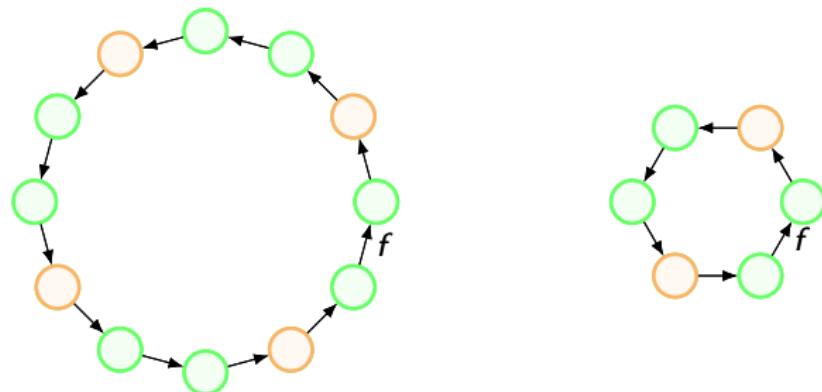
Hellman's algorithm

- ▶ What if y is on a large cycle?



Hellman's algorithm

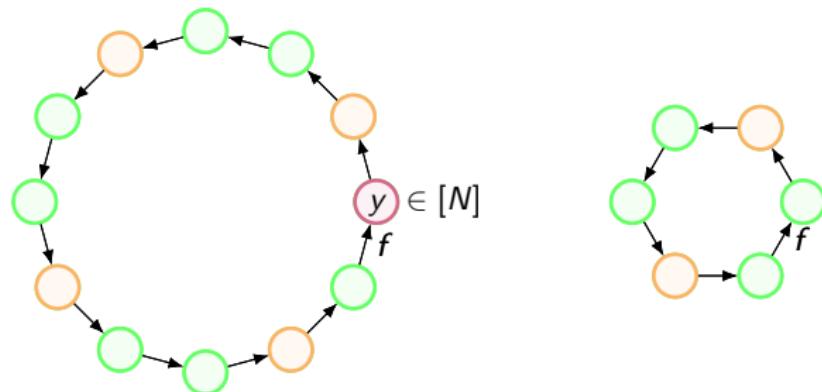
- ▶ What if y is on a large cycle?



- ▶ In a preprocessing step, store points uniformly spaced around each cycle.

Hellman's algorithm

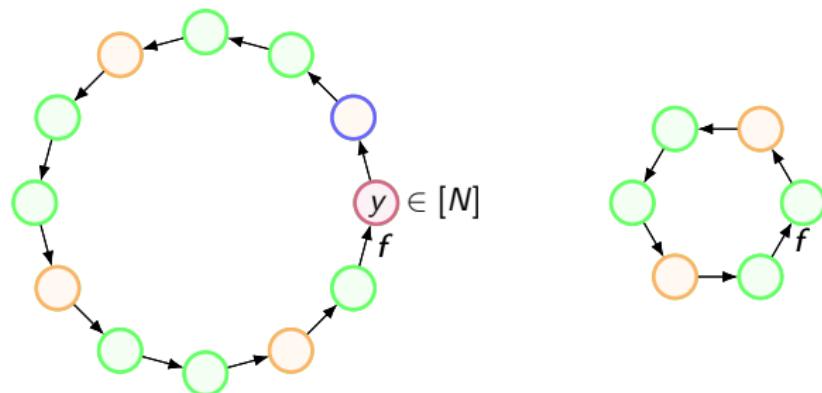
- ▶ What if y is on a large cycle?



- ▶ In a preprocessing step, store points uniformly spaced around each cycle.

Hellman's algorithm

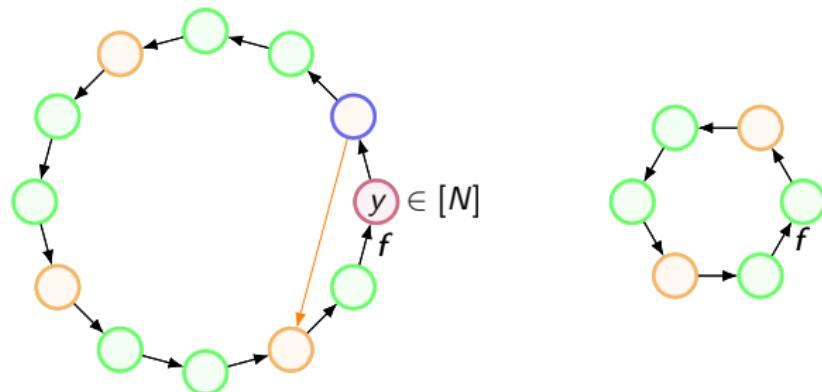
- ▶ What if y is on a large cycle?



- ▶ In a preprocessing step, store points uniformly spaced around each cycle.

Hellman's algorithm

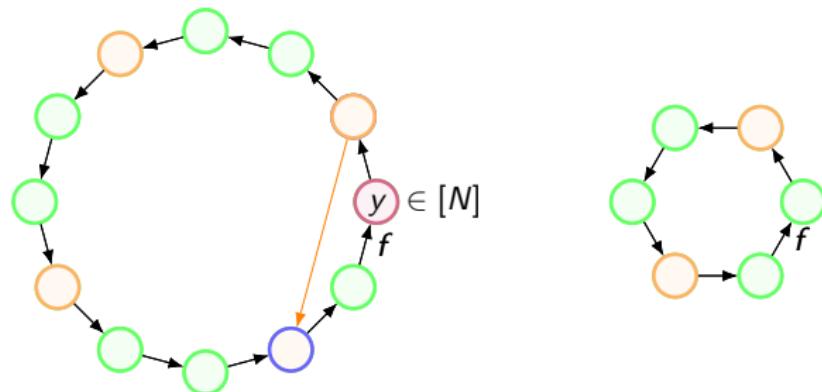
- ▶ What if y is on a large cycle?



- ▶ In a preprocessing step, store points uniformly spaced around each cycle.

Hellman's algorithm

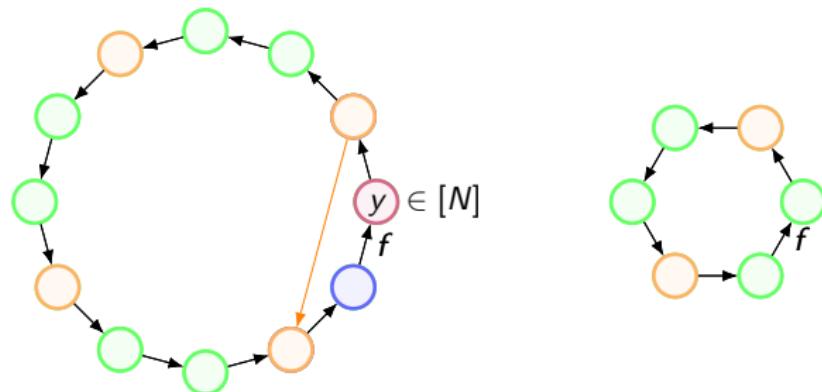
- ▶ What if y is on a large cycle?



- ▶ In a preprocessing step, store points uniformly spaced around each cycle.

Hellman's algorithm

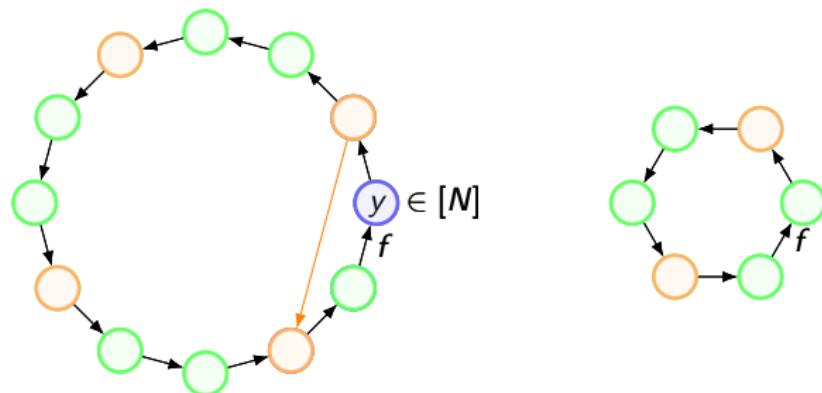
- ▶ What if y is on a large cycle?



- ▶ In a preprocessing step, store points uniformly spaced around each cycle.

Hellman's algorithm

- ▶ What if y is on a large cycle?



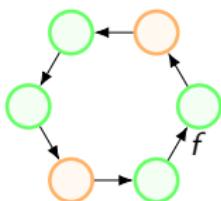
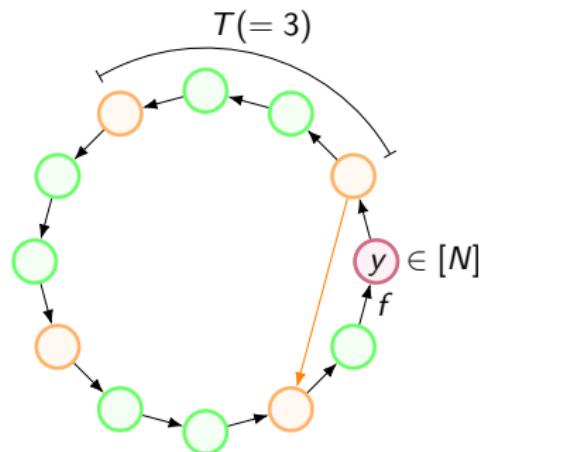
- ▶ In a preprocessing step, store points uniformly spaced around each cycle.

Analysis

- ▶ If we space the stored points T hops apart:

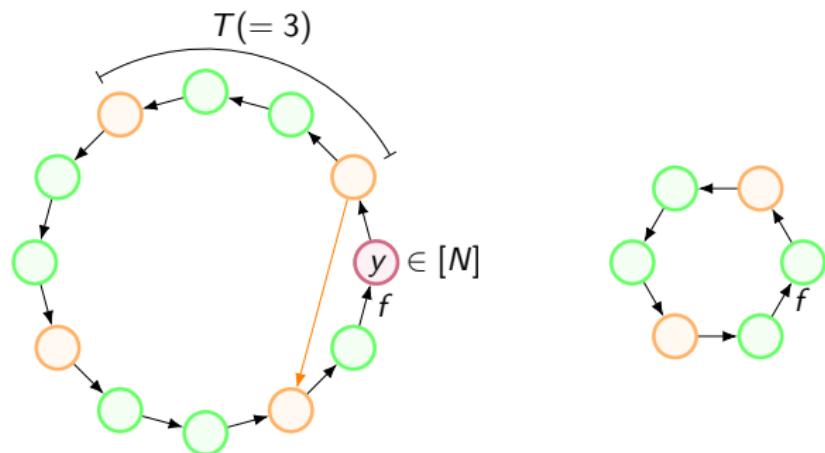
Analysis

- If we space the stored points T hops apart:



Analysis

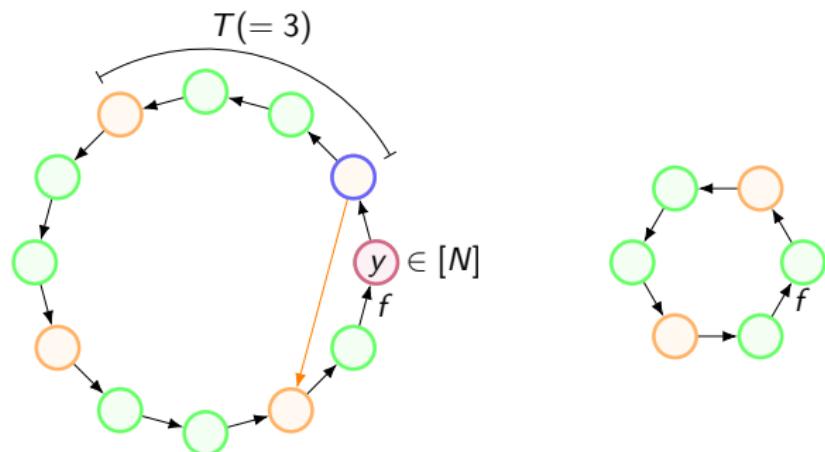
- ▶ If we space the stored points T hops apart:



- ▶ We need T evaluations of f to invert y .

Analysis

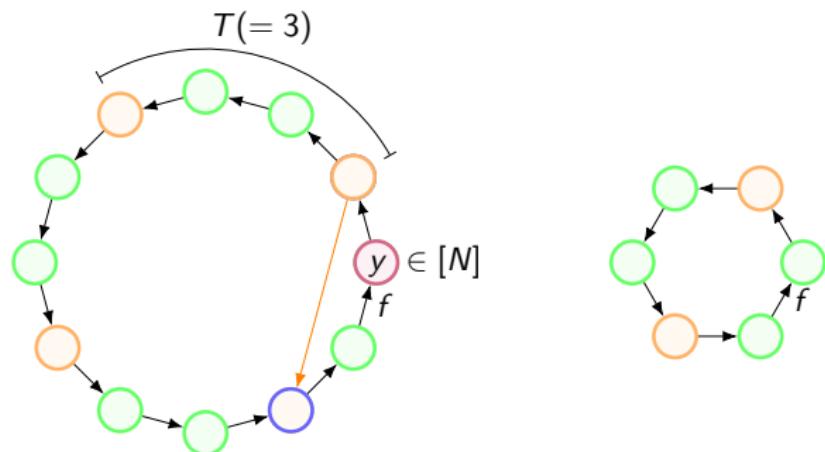
- ▶ If we space the stored points T hops apart:



- ▶ We need T evaluations of f to invert y .

Analysis

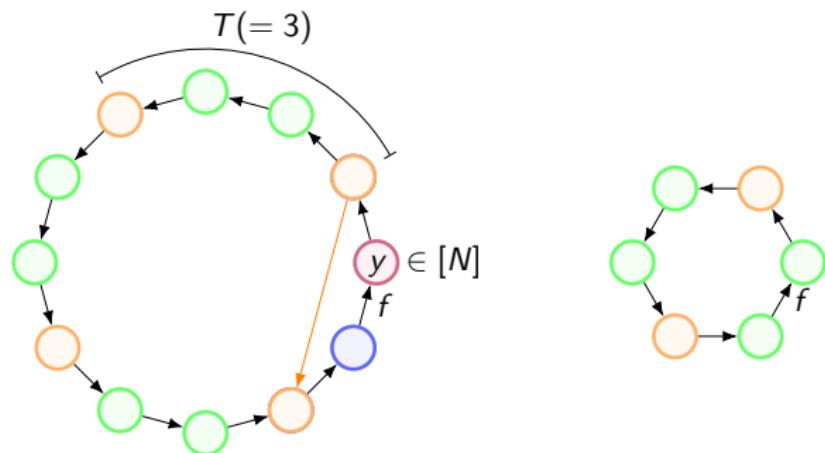
- ▶ If we space the stored points T hops apart:



- ▶ We need T evaluations of f to invert y .

Analysis

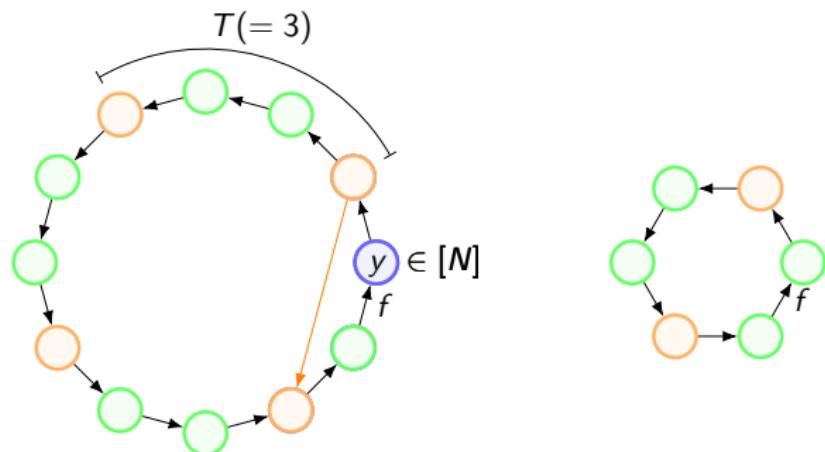
- ▶ If we space the stored points T hops apart:



- ▶ We need T evaluations of f to invert y .

Analysis

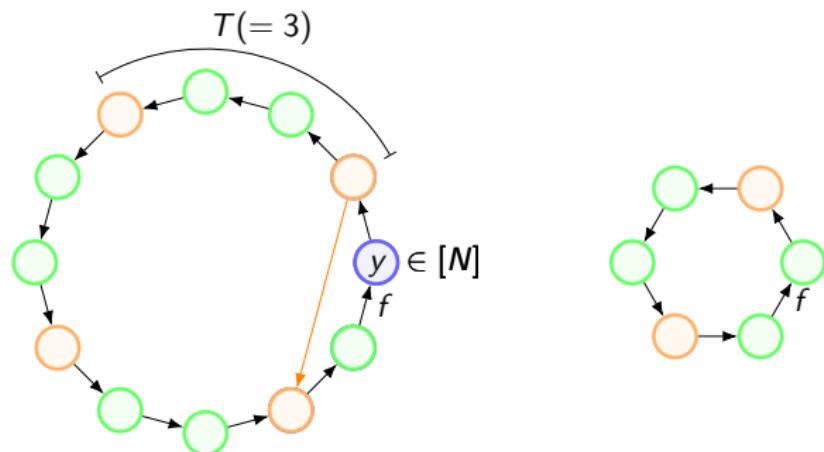
- ▶ If we space the stored points T hops apart:



- ▶ We need T evaluations of f to invert y .

Analysis

- ▶ If we space the stored points T hops apart:



- ▶ We need T evaluations of f to invert y .
- ▶ We need to store about N/T points total.

Stepping back

- ▶ **Goal:** design a pair of algorithms $(\mathcal{P}, \mathcal{A})$ such that for all f and all $y \in f([N])$,

$$\Pr[\alpha \leftarrow \mathcal{P}(f); x \leftarrow \mathcal{A}^f(\alpha, y); f(x) = y] \geq \text{constant}.$$

Stepping back

- ▶ **Goal:** design a pair of algorithms $(\mathcal{P}, \mathcal{A})$ such that for all f and all $y \in f([N])$,

$$\Pr[\alpha \leftarrow \mathcal{P}(f); x \leftarrow \mathcal{A}^f(\alpha, y); f(x) = y] \geq \text{constant}.$$

- ▶ In this model, \mathcal{P} and \mathcal{A} have **unbounded** computational power.

Stepping back

- ▶ **Goal:** design a pair of algorithms $(\mathcal{P}, \mathcal{A})$ such that for all f and all $y \in f([N])$,

$$\Pr[\alpha \leftarrow \mathcal{P}(f); x \leftarrow \mathcal{A}^f(\alpha, y); f(x) = y] \geq \text{constant}.$$

- ▶ In this model, \mathcal{P} and \mathcal{A} have **unbounded** computational power.
- ▶ We aim to **minimize** the bitlength S of α , and the number of queries T that \mathcal{A} makes to f .

Application scenarios

- ▶ The NSA wants to break a widely used cryptographic function, such as AES-128

Application scenarios

- ▶ The NSA wants to break a widely used cryptographic function, such as AES-128
- ▶ Hackers want to recover passwords from a stolen database of password hashes (Rainbow Tables)

Application scenarios

- ▶ The NSA wants to break a widely used cryptographic function, such as AES-128
- ▶ Hackers want to recover passwords from a stolen database of password hashes (Rainbow Tables)
- ▶ Theoretical computer scientists want better algorithms for 3-SUM [GGH⁺20], 

Application scenarios

- ▶ The NSA wants to break a widely used cryptographic function, such as AES-128
- ▶ Hackers want to recover passwords from a stolen database of password hashes (Rainbow Tables)
- ▶ Theoretical computer scientists want better algorithms for 3-SUM [GGH⁺20],

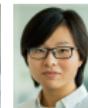


multiparty pointer jumping [CK19],



Application scenarios

- ▶ The NSA wants to break a widely used cryptographic function, such as AES-128
- ▶ Hackers want to recover passwords from a stolen database of password hashes (Rainbow Tables)
- ▶ Theoretical computer scientists want better algorithms for 3-SUM [GGH⁺20],



multiparty pointer jumping [CK19],



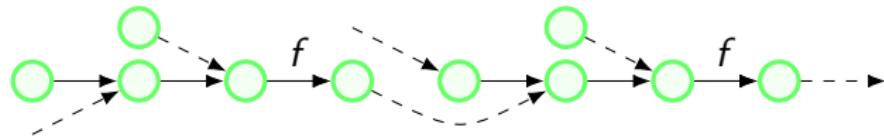
systematic substring search [CK19], ...

Beyond permutations

- ▶ Preprocessing stores the endpoints of disjoint paths.

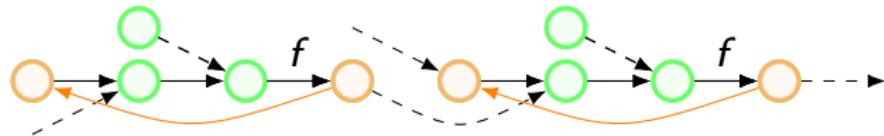
Beyond permutations

- ▶ Preprocessing stores the endpoints of disjoint paths.



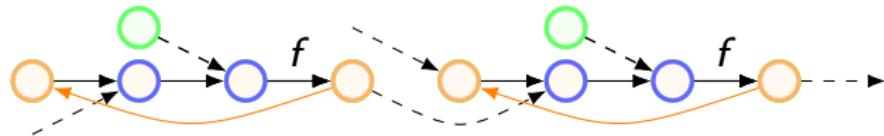
Beyond permutations

- ▶ Preprocessing stores the endpoints of disjoint paths.



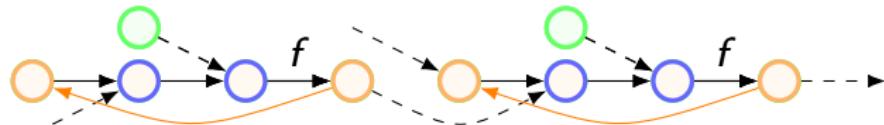
Beyond permutations

- ▶ Preprocessing stores the endpoints of disjoint paths.



Beyond permutations

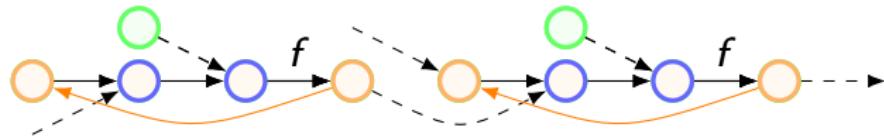
- ▶ Preprocessing stores the endpoints of disjoint paths.



- ▶ No longer possible to cover the entire range.

Beyond permutations

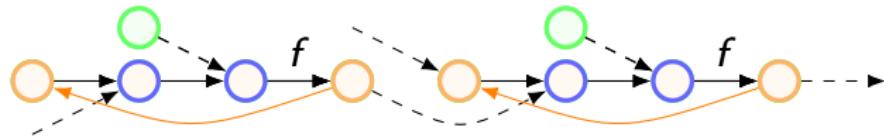
- ▶ Preprocessing stores the endpoints of disjoint paths.



- ▶ No longer possible to cover the entire range.
- ▶ But, can still cover a small fraction with disjoint paths.

Beyond permutations

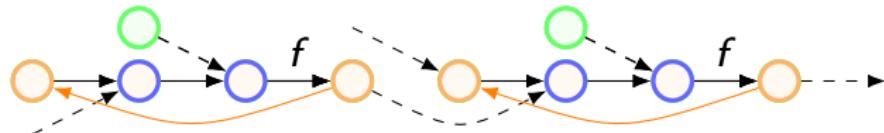
- ▶ Preprocessing stores the endpoints of disjoint paths.



- ▶ No longer possible to cover the entire range.
- ▶ But, can still cover a small fraction with disjoint paths.
- ▶ To boost the coverage, observe that inverting $g \circ f$ on $g(y)$ is often enough to invert f on y .

Beyond permutations

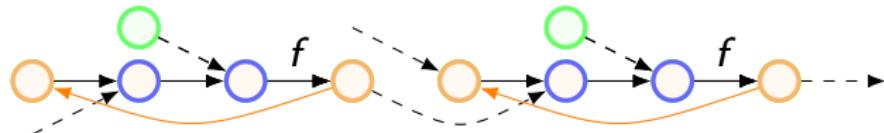
- ▶ Preprocessing stores the endpoints of disjoint paths.



- ▶ No longer possible to cover the entire range.
- ▶ But, can still cover a small fraction with disjoint paths.
- ▶ To boost the coverage, observe that inverting $g \circ f$ on $g(y)$ is often enough to invert f on y .
- ▶ So, can repeatedly apply the basic scheme to many compositions $g_i \circ f$, for suitably chosen “rerandomization” functions g_i .

Beyond permutations

- ▶ Preprocessing stores the endpoints of disjoint paths.

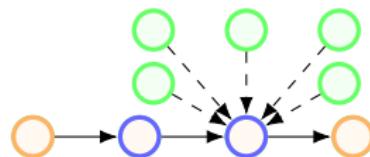


- ▶ No longer possible to cover the entire range.
- ▶ But, can still cover a small fraction with disjoint paths.
- ▶ To boost the coverage, observe that inverting $g \circ f$ on $g(y)$ is often enough to invert f on y .
- ▶ So, can repeatedly apply the basic scheme to many compositions $g_i \circ f$, for suitably chosen “rerandomization” functions g_i .
- ▶ For *random* functions, Hellman showed (heuristically) this can be made to work.



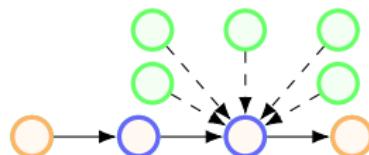
Beyond permutations

- ▶ Hellman's argument fails for arbitrary functions.



Beyond permutations

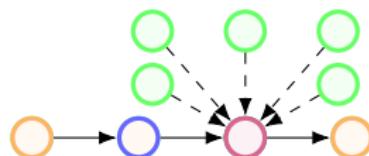
- ▶ Hellman's argument fails for arbitrary functions.



- ▶ Arbitrary functions can have “junction points” with many inverses.

Beyond permutations

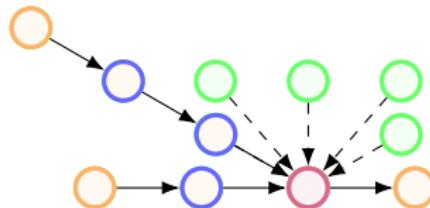
- ▶ Hellman's argument fails for arbitrary functions.



- ▶ Arbitrary functions can have “junction points” with many inverses.

Beyond permutations

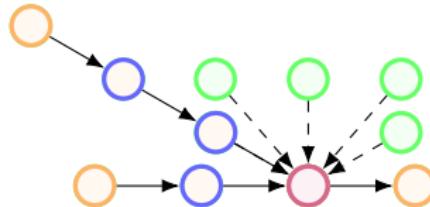
- ▶ Hellman's argument fails for arbitrary functions.



- ▶ Arbitrary functions can have “junction points” with many inverses.

Beyond permutations

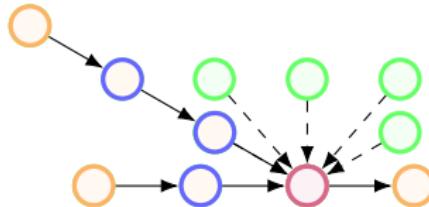
- ▶ Hellman's argument fails for arbitrary functions.



- ▶ Arbitrary functions can have “junction points” with many inverses.
- ▶ Paths collide at these points, causing all sorts of problems.
- ▶ Fiat and Naor deal with this by storing $\alpha = (\alpha', L)$, where L contains junction points along with their inverses.

Beyond permutations

- ▶ Hellman's argument fails for arbitrary functions.



- ▶ Arbitrary functions can have “junction points” with many inverses.
- ▶ Paths collide at these points, causing all sorts of problems.
- ▶ Fiat and Naor deal with this by storing $\alpha = (\alpha', L)$, where L contains junction points along with their inverses.
- ▶ Intuitively, α' is the data structure for a *restriction* of f that avoids the junction points in L .
- ▶ More precisely, the “rerandomization” functions are sampled using rejection sampling so that their range is $[N] - L$.

The Story So Far

Result	Applies To	Tradeoff	Key Point
Hellman 1980 	permutations	$T \lesssim N/S$	$S = T \lesssim \sqrt{N}$

The Story So Far

Result	Applies To	Tradeoff	Key Point
Hellman 1980 	permutations	$T \lesssim N/S$	$S = T \lesssim \sqrt{N}$
Yao 1990 	permutations	$T \gtrsim N/S$	$S = T \gtrsim \sqrt{N}$

The Story So Far

Result	Applies To	Tradeoff	Key Point
Hellman 1980 	permutations	$T \lesssim N/S$	$S = T \lesssim \sqrt{N}$
Yao 1990 	permutations	$T \gtrsim N/S$	$S = T \gtrsim \sqrt{N}$
Hellman 1980 	random f	$T \lesssim N^2/S^2$	$S = T \lesssim N^{2/3}$

The Story So Far

Result	Applies To	Tradeoff	Key Point
Hellman 1980 	permutations	$T \lesssim N/S$	$S = T \lesssim \sqrt{N}$
Yao 1990 	permutations	$T \gtrsim N/S$	$S = T \gtrsim \sqrt{N}$
Hellman 1980 	random f	$T \lesssim N^2/S^2$	$S = T \lesssim N^{2/3}$
Fiat-Naor 1991 	all functions	$T \lesssim N^3/S^3$	$S = T \lesssim N^{3/4}$

The Story So Far

Result	Applies To	Tradeoff	Key Point
Hellman 1980 	permutations	$T \lesssim N/S$	$S = T \lesssim \sqrt{N}$
Yao 1990 	permutations	$T \gtrsim N/S$	$S = T \gtrsim \sqrt{N}$
Hellman 1980 	random f	$T \lesssim N^2/S^2$	$S = T \lesssim N^{2/3}$
Fiat-Naor 1991 	all functions	$T \lesssim N^3/S^3$	$S = T \lesssim N^{3/4}$

- Q: Can we improve Fiat-Naor? Can we improve Yao's lower bound?

The Story So Far

Result	Applies To	Tradeoff	Key Point
Hellman 1980 	permutations	$T \lesssim N/S$	$S = T \lesssim \sqrt{N}$
Yao 1990 	permutations	$T \gtrsim N/S$	$S = T \gtrsim \sqrt{N}$
Hellman 1980 	random f	$T \lesssim N^2/S^2$	$S = T \lesssim N^{2/3}$
Fiat-Naor 1991 	all functions	$T \lesssim N^3/S^3$	$S = T \lesssim N^{3/4}$

- ▶ Q: Can we improve Fiat-Naor? Can we improve Yao's lower bound?
- ▶ A: Sort of and sort of!

Part 1: A Simple Improvement to Fiat-Naor

Result		Applies To	Tradeoff	Key Point
Hellman	1980	permutations	$T \lesssim N/S$	$S = T \lesssim \sqrt{N}$
		permutations	$T \gtrsim N/S$	$S = T \gtrsim \sqrt{N}$
Hellman	1980	random f	$T \lesssim N^2/S^2$	$S = T \lesssim N^{2/3}$
Fiat-Naor	1991	all functions	$T \lesssim N^3/S^3$	$S = T \lesssim N^{3/4}$
				

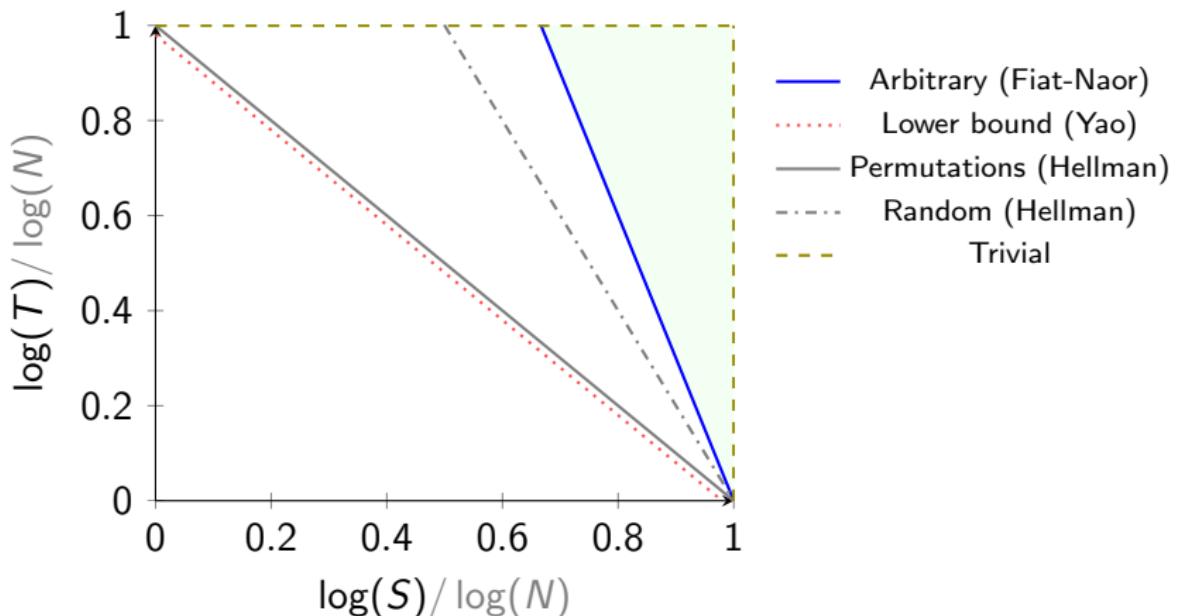
Part 1: A Simple Improvement to Fiat-Naor

Result		Applies To	Tradeoff	Key Point
Hellman 	1980	permutations	$T \lesssim N/S$	$S = T \lesssim \sqrt{N}$
Yao 	1990	permutations	$T \gtrsim N/S$	$S = T \gtrsim \sqrt{N}$
Hellman 	1980	random f	$T \lesssim N^2/S^2$	$S = T \lesssim N^{2/3}$
Fiat-Naor 	1991	all functions	$T \lesssim N^3/S^3$	$S = T \lesssim N^{3/4}$
This work		all functions	$T \lesssim N^3/(S^2 T)$	$S = T \lesssim N^{3/4}$

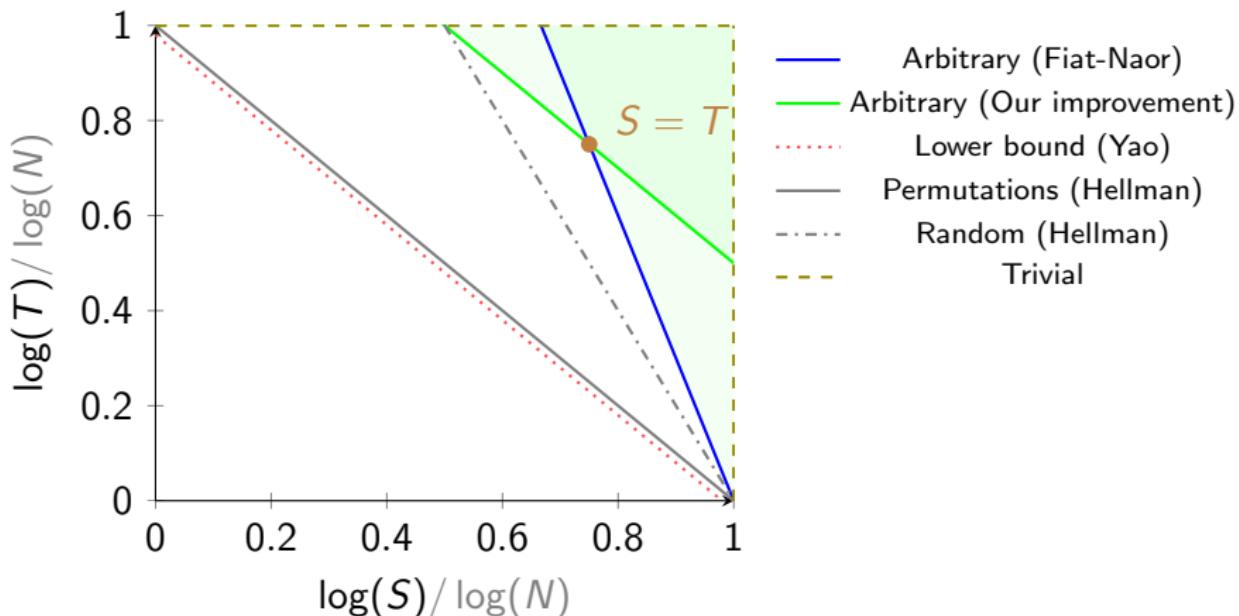
Part 1: A Simple Improvement to Fiat-Naor

Result		Applies To	Tradeoff	Key Point
Hellman 	1980	permutations	$T \lesssim N/S$	$S = T \lesssim \sqrt{N}$
Yao 	1990	permutations	$T \gtrsim N/S$	$S = T \gtrsim \sqrt{N}$
Hellman 	1980	random f	$T \lesssim N^2/S^2$	$S = T \lesssim N^{2/3}$
Fiat-Naor 	1991	all functions	$T \lesssim N^3/S^3$	$S = T \lesssim N^{3/4}$
This work 		all functions	$T \lesssim N^3/(S^2 T)$ $T \lesssim N^{3/2}/S$	$S = T \lesssim N^{3/4}$

A Simple Improvement to Fiat-Naor



A Simple Improvement to Fiat-Naor



A Simple Improvement to Fiat-Naor

- ▶ Recall that Fiat and Naor's preprocessing and online algorithms must agree on a list L of "junction points".

A Simple Improvement to Fiat-Naor

- ▶ Recall that Fiat and Naor's preprocessing and online algorithms must agree on a list L of "junction points".
- ▶ We observe that the tradeoff $T \lesssim N^3/S^3$ comes from

$$T \lesssim \frac{1}{|L|} \cdot \frac{N^3}{S^2}.$$

A Simple Improvement to Fiat-Naor

- ▶ Recall that Fiat and Naor's preprocessing and online algorithms must agree on a list L of "junction points".
- ▶ We observe that the tradeoff $T \lesssim N^3/S^3$ comes from

$$T \lesssim \frac{1}{|L|} \cdot \frac{N^3}{S^2}.$$

- ▶ Fiat and Naor get $|L| \simeq S$, but this is the hard limit, since L needs to fit into S -bit advice α .
- ▶ Or does it?

A Simple Improvement to Fiat-Naor

- ▶ Fiat and Naor's list L actually consists of images $f(x_i)$ of random points $x_i \sim [N]$.

A Simple Improvement to Fiat-Naor

- ▶ Fiat and Naor's list L actually consists of images $f(x_i)$ of random points $x_i \sim [N]$.
- ▶ Our idea: Instead of reading L from α , \mathcal{A} recovers L by evaluating f on the same random points x_i .

A Simple Improvement to Fiat-Naor

- ▶ Fiat and Naor's list L actually consists of images $f(x_i)$ of random points $x_i \sim [N]$.
- ▶ Our idea: Instead of reading L from α , \mathcal{A} recovers L by evaluating f on the same random points x_i .
- ▶ This allows $|L| \simeq T$, so we can get $T \lesssim N^3/(S^2 T)$, or

$$T \lesssim N^{3/2}/S.$$

A Simple Improvement to Fiat-Naor

- ▶ Fiat and Naor's list L actually consists of images $f(x_i)$ of random points $x_i \sim [N]$.
- ▶ Our idea: Instead of reading L from α , \mathcal{A} recovers L by evaluating f on the same random points x_i .
- ▶ This allows $|L| \simeq T$, so we can get $T \lesssim N^3/(S^2 T)$, or

$$T \lesssim N^{3/2}/S.$$

- ▶ But I've cheated here...

A Simple Improvement to Fiat-Naor

- ▶ Fiat and Naor's list L actually consists of images $f(x_i)$ of random points $x_i \sim [N]$.
- ▶ Our idea: Instead of reading L from α , \mathcal{A} recovers L by evaluating f on the same random points x_i .
- ▶ This allows $|L| \simeq T$, so we can get $T \lesssim N^3/(S^2 T)$, or

$$T \lesssim N^{3/2}/S.$$

- ▶ But I've cheated here...
- ▶ How do \mathcal{A} and \mathcal{P} agree on the *same* list of random values x_i ?

Sharing Randomness

- We show that we can assume *shared randomness* without loss of generality.

Sharing Randomness

- ▶ We show that we can assume *shared randomness* without loss of generality.
- ▶ More precisely, shared randomness costs an additive $O(\log N)$ in S .

Sharing Randomness

- ▶ We show that we can assume *shared randomness* without loss of generality.
- ▶ More precisely, shared randomness costs an additive $O(\log N)$ in S .
- ▶ The proof adapts *Newman's lemma* [New91]  from communication complexity.

Sharing Randomness

- ▶ We show that we can assume *shared randomness* without loss of generality.
- ▶ More precisely, shared randomness costs an additive $O(\log N)$ in S .
- ▶ The proof adapts *Newman's lemma* [New91]  from communication complexity.
- ▶ The idea: rather than storing the random bits in full, store a random *index* into a fixed list of strings.

Sharing Randomness

- ▶ We show that we can assume *shared randomness* without loss of generality.
- ▶ More precisely, shared randomness costs an additive $O(\log N)$ in S .
- ▶ The proof adapts *Newman's lemma* [New91]  from communication complexity.
- ▶ The idea: rather than storing the random bits in full, store a random *index* into a fixed list of strings.
- ▶ If the number $\Sigma = O(N)$ of fixed strings is large enough, this will be almost as good as fully random.

Sharing Randomness

- ▶ We show that we can assume *shared randomness* without loss of generality.
- ▶ More precisely, shared randomness costs an additive $O(\log N)$ in S .
- ▶ The proof adapts *Newman's lemma* [New91]  from communication complexity.
- ▶ The idea: rather than storing the random bits in full, store a random *index* into a fixed list of strings.
- ▶ If the number $\Sigma = O(N)$ of fixed strings is large enough, this will be almost as good as fully random.
- ▶ In practice, can instantiate a random oracle.

Part 2: A Lower Bound

- ▶ Recall that Yao's lower bound (for inverting *arbitrary* functions) hasn't been improved in 30+ years.

Part 2: A Lower Bound

- ▶ Recall that Yao's lower bound (for inverting *arbitrary* functions) hasn't been improved in 30+ years.
- ▶ Corrigan-Gibbs and Kogan   [CK19]: any small improvement \implies new lower bounds in circuit complexity.

Part 2: A Lower Bound

- ▶ Recall that Yao's lower bound (for inverting *arbitrary* functions) hasn't been improved in 30+ years.
- ▶ Corrigan-Gibbs and Kogan   [CK19]: any small improvement \implies new lower bounds in circuit complexity.
- ▶ Even improving Yao's bound just for *non-adaptive* algorithms would do it!

Part 2: A Lower Bound

- ▶ Recall that Yao's lower bound (for inverting *arbitrary* functions) hasn't been improved in 30+ years.
- ▶ Corrigan-Gibbs and Kogan   [CK19]: any small improvement \implies new lower bounds in circuit complexity.
- ▶ Even improving Yao's bound just for *non-adaptive* algorithms would do it!
- ▶ \mathcal{A} is non-adaptive if its evaluation points x_1, \dots, x_T are chosen up front, before any evaluations of f are seen.

Part 2: A Lower Bound

- ▶ Recall that Yao's lower bound (for inverting *arbitrary* functions) hasn't been improved in 30+ years.
- ▶ Corrigan-Gibbs and Kogan   [CK19]: any small improvement \implies new lower bounds in circuit complexity.
- ▶ Even improving Yao's bound just for *non-adaptive* algorithms would do it!
- ▶ \mathcal{A} is non-adaptive if its evaluation points x_1, \dots, x_T are chosen up front, before any evaluations of f are seen.
- ▶ Non-adaptive algorithms seem very weak. Hellman's algorithm is *very* adaptive.

Part 2: A Lower Bound

- ▶ Recall that Yao's lower bound (for inverting *arbitrary* functions) hasn't been improved in 30+ years.
- ▶ Corrigan-Gibbs and Kogan   [CK19]: any small improvement \implies new lower bounds in circuit complexity.
- ▶ Even improving Yao's bound just for *non-adaptive* algorithms would do it!
- ▶ \mathcal{A} is non-adaptive if its evaluation points x_1, \dots, x_T are chosen up front, before any evaluations of f are seen.
- ▶ Non-adaptive algorithms seem very weak. Hellman's algorithm is *very* adaptive.
- ▶ Corrigan-Gibbs and Kogan speculated that there is no non-adaptive algorithm with

$$S = o(N \log N) \text{ and } T = o(N).$$

A Simple Non-Adaptive Algorithm

- ▶ Corrigan-Gibbs and Kogan speculated that there is no non-adaptive algorithm with

$$S = o(N \log N) \text{ and } T = o(N).$$

- ▶ We observe that there IS in fact a very simple algorithm, that (barely!) outperforms the trivial inverter.

A Simple Non-Adaptive Algorithm

- ▶ Corrigan-Gibbs and Kogan speculated that there is no non-adaptive algorithm with

$$S = o(N \log N) \text{ and } T = o(N).$$

- ▶ We observe that there IS in fact a very simple algorithm, that (barely!) outperforms the trivial inverter.
- ▶ For each range element, just store all but $\log(T)$ bits of one of its inverses (implicitly defining T candidate inverses). This achieves $S = O(N \log(N/T))$.

A Simple Non-Adaptive Algorithm

- ▶ Corrigan-Gibbs and Kogan speculated that there is no non-adaptive algorithm with

$$S = o(N \log N) \text{ and } T = o(N).$$

- ▶ We observe that there IS in fact a very simple algorithm, that (barely!) outperforms the trivial inverter.
- ▶ For each range element, just store all but $\log(T)$ bits of one of its inverses (implicitly defining T candidate inverses). This achieves $S = O(N \log(N/T))$.

A Matching Lower Bound

- ▶ The very simple algorithm is a *guess-and-check algorithm*. That is, it can be thought of as making *no* queries to f , and returning T candidate inverses, one of which is guaranteed to be correct.

A Matching Lower Bound

- ▶ The very simple algorithm is a *guess-and-check algorithm*. That is, it can be thought of as making *no* queries to f , and returning T candidate inverses, one of which is guaranteed to be correct.
- ▶ We show that all *guess-and-check* algorithms satisfy the matching lower bound $S = \Omega(N \log(N/T))$.

A Matching Lower Bound

- ▶ For simplicity, assume a guess-and-check algorithm that always succeeds, with parameters S, T . We use a compression argument.

A Matching Lower Bound

- ▶ For simplicity, assume a guess-and-check algorithm that always succeeds, with parameters S, T . We use a compression argument.
- ▶ Theorem: Can encode any permutation f using $S + N \log T$ bits.

A Matching Lower Bound

- ▶ For simplicity, assume a guess-and-check algorithm that always succeeds, with parameters S, T . We use a compression argument.
- ▶ Theorem: Can encode any permutation f using $S + N \log T$ bits.

$$\implies S + N \log T = \Omega(N \log N)$$

$$\implies S = \Omega(N \log(N/T)).$$

A Matching Lower Bound

- ▶ For simplicity, assume a guess-and-check algorithm that always succeeds, with parameters S, T . We use a compression argument.
- ▶ Theorem: Can encode any permutation f using $S + N \log T$ bits.

$$\begin{aligned}\implies S + N \log T &= \Omega(N \log N) \\ \implies S &= \Omega(N \log(N/T)).\end{aligned}$$

- ▶ Proof:
 - ▶ Encoder computes $\alpha \leftarrow \mathcal{P}(f)$.

A Matching Lower Bound

- ▶ For simplicity, assume a guess-and-check algorithm that always succeeds, with parameters S, T . We use a compression argument.
- ▶ Theorem: Can encode any permutation f using $S + N \log T$ bits.

$$\begin{aligned}\implies S + N \log T &= \Omega(N \log N) \\ \implies S &= \Omega(N \log(N/T)).\end{aligned}$$

- ▶ Proof:
 - ▶ Encoder computes $\alpha \leftarrow \mathcal{P}(f)$.
 - ▶ For each $y \in [N]$, encoder runs $\mathcal{A}(\alpha, y)$ and receives x_1, \dots, x_T . It writes down the $i_y \in [T]$ that satisfies $f(x_{i_y}) = y$.

A Matching Lower Bound

- ▶ For simplicity, assume a guess-and-check algorithm that always succeeds, with parameters S, T . We use a compression argument.
- ▶ Theorem: Can encode any permutation f using $S + N \log T$ bits.

$$\begin{aligned}\implies S + N \log T &= \Omega(N \log N) \\ \implies S &= \Omega(N \log(N/T)).\end{aligned}$$

- ▶ Proof:
 - ▶ Encoder computes $\alpha \leftarrow \mathcal{P}(f)$.
 - ▶ For each $y \in [N]$, encoder runs $\mathcal{A}(\alpha, y)$ and receives x_1, \dots, x_T . It writes down the $i_y \in [T]$ that satisfies $f(x_{i_y}) = y$.
 - ▶ Encoding is $(\alpha, i_1, \dots, i_N)$.

A Matching Lower Bound

- ▶ For simplicity, assume a guess-and-check algorithm that always succeeds, with parameters S, T . We use a compression argument.
- ▶ Theorem: Can encode any permutation f using $S + N \log T$ bits.

$$\begin{aligned}\implies S + N \log T &= \Omega(N \log N) \\ \implies S &= \Omega(N \log(N/T)).\end{aligned}$$

- ▶ Proof:
 - ▶ Encoder computes $\alpha \leftarrow \mathcal{P}(f)$.
 - ▶ For each $y \in [N]$, encoder runs $\mathcal{A}(\alpha, y)$ and receives x_1, \dots, x_T . It writes down the $i_y \in [T]$ that satisfies $f(x_{i_y}) = y$.
 - ▶ Encoding is $(\alpha, i_1, \dots, i_N)$.
 - ▶ For each y , decoder again runs $\mathcal{A}(\alpha, y)$ and receives x_1, \dots, x_T . It sets $f^{-1}(y) = x_{i_y}$.

Function Inversion Revisited: Our Results

- ▶ Improving Fiat-Naor: a

$$T \lesssim N^{3/2}/S$$

algorithm for inverting arbitrary functions.

Function Inversion Revisited: Our Results

- ▶ Improving Fiat-Naor: a

$$T \lesssim N^{3/2}/S$$

algorithm for inverting arbitrary functions.

- ▶ Shared randomness is essentially free

Function Inversion Revisited: Our Results

- ▶ Improving Fiat-Naor: a

$$T \lesssim N^{3/2}/S$$

algorithm for inverting arbitrary functions.

- ▶ Shared randomness is essentially free
- ▶ We observe there is a very simple nontrivial non-adaptive function inversion algorithm, and show it is best possible among *guess-and-check* algorithms.

Function Inversion Revisited: Our Results

- ▶ Improving Fiat-Naor: a

$$T \lesssim N^{3/2}/S$$

algorithm for inverting arbitrary functions.

- ▶ Shared randomness is essentially free
- ▶ We observe there is a very simple nontrivial non-adaptive function inversion algorithm, and show it is best possible among *guess-and-check* algorithms.
- ▶ Not in this talk:
 - ▶ Search-to-decision reductions

Function Inversion Revisited: Our Results

- ▶ Improving Fiat-Naor: a

$$T \lesssim N^{3/2}/S$$

algorithm for inverting arbitrary functions.

- ▶ Shared randomness is essentially free
- ▶ We observe there is a very simple nontrivial non-adaptive function inversion algorithm, and show it is best possible among *guess-and-check* algorithms.
- ▶ Not in this talk:
 - ▶ Search-to-decision reductions
 - ▶ A larger codomain does not make the problem harder (us and [CK19])

Open Problems

- ▶ Only one real open problem—

Open Problems

- ▶ Only one real open problem—**close the gap!**

Open Problems

- ▶ Only one real open problem—**close the gap!**
- ▶ Moonshots:
 - ▶ Improve Yao's lower bound against (general) non-adaptive algorithms?

Open Problems

- ▶ Only one real open problem—**close the gap!**
- ▶ Moonshots:
 - ▶ Improve Yao's lower bound against (general) non-adaptive algorithms?
 - ▶ $S^2 T = N^2$ algorithm for worst-case function inversion?

Open Problems

- ▶ Only one real open problem—**close the gap!**
- ▶ Moonshots:
 - ▶ Improve Yao's lower bound against (general) non-adaptive algorithms?
 - ▶ $S^2 T = N^2$ algorithm for worst-case function inversion?
- ▶ Possibly more tractable:
 - ▶ Better algorithms for inverting a small fraction of the range?
That is, improving on De, Trevisan and Tulsiani [DTT10]?

Open Problems

- ▶ Only one real open problem—**close the gap!**
- ▶ Moonshots:
 - ▶ Improve Yao's lower bound against (general) non-adaptive algorithms?
 - ▶ $S^2 T = N^2$ algorithm for worst-case function inversion?
- ▶ Possibly more tractable:
 - ▶ Better algorithms for inverting a small fraction of the range?
That is, improving on De, Trevisan and Tulsiani [DTT10]?
 - ▶ Better algorithms for injective functions?

Thank you!

I'm happy to take additional questions offline. You can ping me at speters@cs.cornell.edu.

References I

-  Henry Corrigan-Gibbs and Dmitry Kogan.
The function-inversion problem: Barriers and opportunities.
In *TCC*, 2019.
-  Anindya De, Luca Trevisan, and Madhur Tulsiani.
Time space tradeoffs for attacks against one-way functions
and PRGs.
In *CRYPTO*, 2010.
-  Alexander Golovnev, Siyao Guo, Thibaut Horel, Sunoo Park,
and Vinod Vaikuntanathan.
Data structures meet cryptography: 3SUM with preprocessing.

In *STOC*, 2020.

References II

-  Alexander Golovnev, Siyao Guo, Spencer Peters, and Noah Stephens-Davidowitz.
Revisiting time-space tradeoffs for function inversion.
Available at
<https://eccc.weizmann.ac.il/report/2022/145/>, 2022.
-  Martin Hellman.
A cryptanalytic time-memory trade-off.
IEEE Transactions on Information Theory, 26(4):401–406, 1980.
-  Ilan Newman.
Private vs. common random bits in communication complexity.
Information processing letters, 39(2):67–71, 1991.